

211112262  
VAIBHAV PATEL

**ML LAB1**

```
[1]: import numpy as  
      np import pandas  
      as pd  
      import matplotlib.pyplot as plt  
      import math
```

```
[5]: train_dirs=[]  
      test_dirs=[]  
      headers=[0, 0, 0, 0, 0]  
      for dir in os.listdir("./"):  
          if(dir.find("5-fold")!= -1):  
              train_dirs.append("./"+dir+"/train/")  
              test_dirs.append("./"+dir+"/test/")
```

**0.0.1 Linear Regression**

```
[1]: def LinearRegression(train_file, test_file, header):  
      train_df = pd.read_csv(train_file, header=header, delimiter=",")  
      test_df = pd.read_csv(test_file, header=header, delimiter=",")  
  
      X_train = train_df.iloc[:, :-1].values  
      y_train = train_df.iloc[:, -1].values  
  
      X_test = test_df.iloc[:, :-1].values  
      y_test = test_df.iloc[:, -1].values  
  
      from sklearn.linear_model import LinearRegression  
      regressor = LinearRegression()  
  
      regressor.fit(X_train, y_train)  
  
      y_pred =  
  
      regressor.predict(X_test)  
  
      from sklearn.metrics import mean_squared_error, mean_absolute_error, _  
      r2_score  
      mse = mean_squared_error(y_test, y_pred)
```

```

mae = mean_absolute_error(y_test, y_pred)
r2score = r2_score(y_test, y_pred)
return np.array([mse, mae, r2score])

```

```

[7]: for train_dir, test_dir, header in zip(train_dirs, test_dirs, headers):
    train_files=os.listdir(train_dir)
    test_files=os.listdir(test_dir)
    val=np.zeros(3)
    for train, test in zip(train_files, test_files):
        val+=LinearRegression(train_dir+train, test_dir+test, header)
    print(train_dir)
    val/=len(train_files)
    val[0]=math.sqrt(val[0])
    val=pd.DataFrame(val, index=["RMSE", "MSE", "R2"], columns=["Values"])
    print(val)
    print("

```

-----

./concrete-5-fold/train/

	Values
RMSE	10.441372
MSE	8.272734
R2	0.603040

-----

./diabetes-5-fold/train/

	Values
RMSE	0.63949
8	
MSE	0.501970
R2	-0.000552

-----

./mortgage-5-fold/train/

	Values
RMSE	0.12182
4	
MSE	0.083536
R2	0.998404

-----

./plastic-5-fold/train/

	Values
RMSE	1.531465
MSE	1.232442
R2	0.798437

-----

./stock-5-fold/train/

	Values
--	--------

```
RMSE 2.346750
MSE 1.832574
R2 0.870620
-----
```

### 0.0.2 Polynomial Regression of degree 2 and 3

```
[8]: def PolynomialRegression(train_file, test_file, header, degree):
    train_df = pd.read_csv(train_file, header=header, delimiter=";")
    test_df = pd.read_csv(test_file, header=header, delimiter=";")

    X_train = train_df.iloc[:, :-1].values
    y_train = train_df.iloc[:, -1].values

    X_test = test_df.iloc[:, :-1].values
    y_test = test_df.iloc[:, -1].values

    from sklearn.linear_model import LinearRegression
    from sklearn.preprocessing import PolynomialFeatures

    poly_reg=PolynomialFeatures(degree=degree)
    X_poly=poly_reg.fit_transform(X_train)

    regressor = LinearRegression()
    regressor.fit(X_poly, y_train)

    y_pred = regressor.predict(poly_reg.transform(X_test))

    from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

    mse = mean_squared_error(y_test, y_pred)
    mae = mean_absolute_error(y_test, y_pred)
    r2score = r2_score(y_test, y_pred)
    return np.array([mse, mae, r2score])
```

### [9]: Polynomial Regression of degree 2

```
for train_dir, test_dir, header in zip(train_dirs, test_dirs, headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(3)
    for train, test in zip(train_files, test_files):
        val += PolynomialRegression(train_dir + train, test_dir + test, header,
2)
```

```

print(train_dir)
val /= len(train_files)
val[0] =
math.sqrt(val[0])
val = pd.DataFrame(val, index=["RMSE", "MSE", "R2"], columns=["Values"])
print(val)
print("\n")

```

./concrete-5-fold/train/

Values

RMSE 7.66514

3

MSE 5.892674

R2 0.786115

---

./diabetes-5-fold/train/

Values

RMSE 0.56129

7

MSE 0.456880

R2 0.226230

---

./mortgage-5-fold/train/

Values

RMSE 0.12061

6

MSE 0.055462

R2 0.998544

---

./plastic-5-fold/train/

Values

RMSE 1.528545

MSE 1.226209

R2 0.799254

---

./stock-5-fold/train/

Values

RMSE 1.127098

MSE 0.901557

R2 0.970154

---

### Polynomial Regression of degree 3

```

[10]: for train_dir, test_dir, header in zip(train_dirs, test_dirs, headers):
      train_files = os.listdir(train_dir)
      test_files = os.listdir(test_dir)

```

```

val = np.zeros(3)
for train, test in zip(train_files, test_files):
    val += PolynomialRegression(train_dir + train, test_dir + test, header,
3)
    print(train_dir)
    val /= len(train_files)
    val[0] =
    math.sqrt(val[0])
    val = pd.DataFrame(val, index=["RMSE", "MSE", "R2"], columns=["Values"])
    print(val)
    print("-----\n")

```

./concrete-5-fold/train/

Values

RMSE

6.72537

9

MSE 4.674355

R2 0.834883

-----

./diabetes-5-fold/train/

Values

RMSE

0.83851

1

MSE 0.620331

R2 -0.519498

-----

./mortgage-5-fold/train/

Values

RMSE

2.59757

3

MSE 0.512271

R2 0.338858

-----

./plastic-5-fold/train/

Values

RMSE 1.473863

MSE 1.166224

R2 0.813267

-----

./stock-5-fold/train/

Values

RMSE 0.880364

MSE 0.661540

R2 0.981784

-----

### 0.0.3 Regularization in Linear Regression

```
[24]: def Regularization(train_file, test_file, header):
    train_df = pd.read_csv(train_file, header=header, delimiter=",")
    test_df = pd.read_csv(test_file, header=header, delimiter=",")

    X_train = train_df.iloc[:, :-1].values
    y_train = train_df.iloc[:, -1].values

    X_test = test_df.iloc[:, :-1].values
    y_test = test_df.iloc[:, -1].values

    from sklearn.linear_model import Ridge
    alphas = np.array([2**i for i in range(-18, 51, 2)])

    best_mse, best_alpha=float('inf'), None

    from sklearn.metrics import mean_squared_error
    for alpha in alphas: regressor=Ridge(alpha=alpha)
        regressor.fit(X_train, y_train)
        y_pred=regressor.predict(X_test)
        mse=mean_squared_error(y_test, y_pred)

        if mse<best_mse:
            best_mse, best_alpha=mse, alpha
    return np.array([best_mse, best_alpha])

[25]: for train_dir, test_dir, header in zip(train_dirs, test_dirs, headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(2)
    for train, test in zip(train_files, test_files):
        val = Regularization(train_dir + train, test_dir + test, header)
    print(train_dir)
    val[0]=math.sqrt(val[0])
    val = pd.DataFrame(val, index=["Best RMSE", "Best Alpha"],_
columns=["Values"])
    print(val)
    print("-----\n")

# generate a table train/
Values
Best RMSE    10.222187
Best Alpha  1024.000000
-----
```

./diabetes-5-fold/train/  
Values

Best RMSE 0.63437

1

Best Alpha 0.000004

-----

./mortgage-5-fold/train/  
Values

Best RMSE 0.11068

4

Best Alpha 1.000000

-----

./plastic-5-fold/train/  
Values

Best RMSE 1.51057

9

Best Alpha 0.000004

-----

./stock-5-fold/train/  
Values

Best RMSE 2.31511

7

Best Alpha 0.000004

-----

#### 0.0.4 Ridge on Linear Reg for Best MAE

```
[26]: def Regularization(train_file, test_file, header):  
    train_df = pd.read_csv(train_file, header=header, delimiter=",")  
    test_df = pd.read_csv(test_file, header=header, delimiter=",")  
  
    X_train = train_df.iloc[:, :-1].values  
    y_train = train_df.iloc[:, -1].values  
  
    X_test = test_df.iloc[:, :-1].values  
    y_test = test_df.iloc[:, -1].values  
  
    from sklearn.linear_model import Ridge  
    alphas = np.array([2**i for i in range(-18, 51, 2)])  
  
    best_mse, best_alpha=float('inf'), None  
  
    from sklearn.metrics import mean_absolute_error  
    for alpha in alphas:  
        regressor=Ridge(alpha=alpha)  
        regressor.fit(X_train, y_train)  
        y_pred=regressor.predict(X_test)
```



```

mse=mean_absolute_error(y_test, y_pred)

if mse<best_mse:
    best_mse, best_alpha=mse, alpha
return np.array([best_mse, best_alpha])

```

```

[27]: for train_dir, test_dir, header in zip(train_dirs, test_dirs, headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(2)
    for train, test in zip(train_files, test_files):
        val = Regularization(train_dir + train, test_dir + test, header)
    print(train_dir)
    val = pd.DataFrame(val, index=["Best MAE", "Best Alpha"],
columns=["Values"])
    print(val)
    print("-----\n")

```

```

./concrete-5-fold/train/
      Values
Best MAE      8.084741
Best Alpha 4096.000000
-----

```

```

./diabetes-5-fold/train/
      Values
Best MAE    0.47644
Best Alpha  4.00000
-----

```

```

./mortgage-5-fold/train/
      Values
Best MAE
0.07921
9
Best Alpha  1.000000
-----

```

```

./plastic-5-fold/train/
      Values
Best MAE    1.234517
Best Alpha 64.000000
-----

```

```

./stock-5-fold/train/
      Values
Best MAE    1.816376
Best Alpha 256.000000

```

### 0.0.5 Ridge for Linear Reg, Best R2 score

```
[28]: def Regularization(train_file, test_file, header):
    train_df = pd.read_csv(train_file, header=header, delimiter=",")
    test_df = pd.read_csv(test_file, header=header, delimiter=",")

    X_train = train_df.iloc[:, :-1].values
    y_train = train_df.iloc[:, -1].values

    X_test = test_df.iloc[:, :-1].values
    y_test = test_df.iloc[:, -1].values

    from sklearn.linear_model import Ridge
    alphas = np.array([2**i for i in range(-18, 51, 2)])

    best_mse, best_alpha=-float('inf'), None

    from sklearn.metrics import r2_score
    for alpha in alphas:
        regressor=Ridge(alpha=alpha)
        regressor.fit(X_train, y_train)
        y_pred=regressor.predict(X_test)
        mse=r2_score(y_test, y_pred)

        if mse>best_mse:
            best_mse, best_alpha=mse, alpha

[29]: for train_dir, test_dir, header in zip(train_dirs, test_dirs, headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(2)
    for train, test in zip(train_files, test_files):
        val = Regularization(train_dir + train, test_dir + test, header)
    print(train_dir)
    val = pd.DataFrame(val, index=["Best R2 Score", "Best Alpha"],
        columns=["Values"])
    print(val)
    print("-----\n")

# Complete Jupyter train/
      Values
Best R2 Score 0.674538
Best Alpha
1024.000000
```

```

-----
./diabetes-5-fold/train/
                Values
Best      R2      Score
0.319099 Best Alpha
0.000004
-----

```

```

-----
./mortgage-5-fold/train/
                Values
Best      R2      Score
0.998849 Best Alpha
1.000000
-----

```

```

-----
./plastic-5-fold/train/
                Values
Best      R2      Score
0.790441 Best Alpha
0.000004
-----

```

```

-----
./stock-5-fold/train/
                Values
Best      R2      Score
0.866915 Best Alpha
0.000004
-----

```

```
[17]: import warnings
      warnings.filterwarnings('ignore')
```

### 0.0.6 Ridge for All using MSE

```
[30]: def PolynomialRidge(train_file, test_file, header, degree):
      train_df = pd.read_csv(train_file, header=header, delimiter=",")
      test_df = pd.read_csv(test_file, header=header, delimiter=",")

      X_train = train_df.iloc[:, :-1].values
      y_train = train_df.iloc[:, -1].values

      X_test = test_df.iloc[:, :-1].values
      y_test = test_df.iloc[:, -1].values

      from sklearn.linear_model import Ridge
      from sklearn.preprocessing import PolynomialFeatures

      poly_reg=PolynomialFeatures(degree=degree)
      X_poly=poly_reg.fit_transform(X_train)
```

```

alphas=np.array([2**i for i in range(-18, 30)])

best_alpha, best_mse=None, float('inf')
for alpha in alphas:
    regressor = Ridge(alpha=alpha)
    regressor.fit(X_poly, y_train)
    y_pred = regressor.predict(poly_reg.transform(X_test))

    from sklearn.metrics import mean_squared_error,
    mean_absolute_error, r2_score
    mse = mean_squared_error(y_test, y_pred)
    if(mse<best_mse):
        best_mse, best_alpha=mse, alpha

```

```

[31]: for train_dir, test_dir, header in zip(train_dirs, test_dirs, headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(2)
    for train, test in zip(train_files, test_files):
        val += PolynomialRidge(train_dir + train, test_dir + test, header, 2)
    print(train_dir)
    val /= len(train_files)
    val[0] =
    math.sqrt(val[0])
    val = pd.DataFrame(val, index=["RMSE", "Alpha"], columns=["Values"])
    print(val)
    print("-----\n")

```

```

./concrete-5-fold/train/
Values

```

```

RMSE
7.646834e+00
Alpha 3.355445e+06
-----

```

```

./diabetes-5-fold/train/

```

```

Values
RMSE 5.275634e-
01 Alpha
1.677722e+06
-----

```

```

./mortgage-5-fold/train/

```

```

Values
RMSE 0.075369
Alpha 153.900001

```

```

-----
./plastic-5-fold/train/
      Values
RMSE      1.527185
Alpha 26265.600002
-----

```

```

-----
./stock-5-fold/train/
      Values
RMSE
      1.11370
6
Alpha 5.200002
-----

```

```

[21]: for train_dir, test_dir, header in zip(train_dirs, test_dirs, headers):
      train_files = os.listdir(train_dir)
      test_files = os.listdir(test_dir)

      val = np.zeros(2)
      for train, test in zip(train_files, test_files):
          val += PolynomialRidge(train_dir + train, test_dir + test, header, 3)
      print(train_dir)
      val /= len(train_files)
      val[0] =
      math.sqrt(val[0])
      val = pd.DataFrame(val, index=["RMSE", "Alpha"], columns=["Values"])
      print(val)
      print("-----\n")

```

```

-----
./concrete-5-fold/train/
      Values
RMSE
6.089751e+00
Alpha 1.074266e+08
-----

```

```

-----
./diabetes-5-fold/train/
      Values
RMSE 6.234885e-
01      Alpha
1.090521e+08
-----

```

```

-----
./mortgage-5-fold/train/
      Values
RMSE 6.824711e-
02      Alpha
1.250165e+08
-----

```

```

-----
./plastic-5-fold/train/

```

	Values
RMSE	1.472243
Alpha	25.625002

-----

	Values
RMSE	0.836642
Alpha	16896.000001

-----

### 0.0.7 Ridge for all using MAE

```
[22]: def PolynomialRidge(train_file, test_file, header, degree):
    train_df = pd.read_csv(train_file, header=header, delimiter=",")
    test_df = pd.read_csv(test_file, header=header, delimiter=",")

    X_train = train_df.iloc[:, :-1].values
    y_train = train_df.iloc[:, -1].values

    X_test = test_df.iloc[:, :-1].values
    y_test = test_df.iloc[:, -1].values

    from sklearn.linear_model import Ridge
    from sklearn.preprocessing import PolynomialFeatures

    poly_reg=PolynomialFeatures(degree=degree)
    X_poly=poly_reg.fit_transform(X_train)

    alphas=np.array([2**i for i in range(-18, 30)])

    best_alpha, best_mse=None, float('inf')
    for alpha in alphas:
        regressor = Ridge(alpha=alpha)
        regressor.fit(X_poly, y_train)
        y_pred = regressor.predict(poly_reg.transform(X_test))

        from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
        mse = mean_absolute_error(y_test, y_pred)
        if(mse<best_mse):
            best_mse, best_alpha=mse, alpha

    return np.array([best_mse, best_alpha])
```

```
[23]: for train_dir, test_dir, header in zip(train_dirs, test_dirs, headers):
      train_files = os.listdir(train_dir)
      test_files = os.listdir(test_dir)

      val = np.zeros(2)
      for train, test in zip(train_files, test_files):
          val += PolynomialRidge(train_dir + train, test_dir + test, header, 3)
      print(train_dir)
      val /= len(train_files)
      val = pd.DataFrame(val, index=["MAE", "Alpha"], columns=["Values"])
      print(val)
      print("-----\n")
```

```
./concrete-5-fold/train/
      Values
```

```
MAE      4.500617e+0
```

```
Alpha    1.074791e+0
```

```
8
```

```
-----
```

```
./diabetes-5-fold/train/
      Values
```

```
MAE      4.660749e-01
```

```
Alpha    1.078460e+0
```

```
8
```

```
-----
```

```
./mortgage-5-fold/train/
      Values
```

```
MAE      4.360656e-02
```

```
Alpha    1.183318e+0
```

```
8
```

```
-----
```

```
./plastic-5-fold/train/
      Values
```

```
MAE      1.164341
```

```
Alpha    0.362501
```

```
-----
```

```
./stock-5-fold/train/
      Values
```

```
MAE      0.647212
```

```
Alpha    7782.400196
```

```
-----
```

```
[32]: for train_dir, test_dir, header in zip(train_dirs, test_dirs, headers):
      train_files = os.listdir(train_dir)
      test_files = os.listdir(test_dir)
```

```

val = np.zeros(2)
for train, test in zip(train_files, test_files):
    val += PolynomialRidge(train_dir + train, test_dir + test, header, 2)
print(train_dir)
val /= len(train_files)
val = pd.DataFrame(val, index=["MAE", "Alpha"], columns=["Values"])
print(val)
print("-----\n")

```

```

./concrete-5-fold/train/
      Values

```

```

MAE      5.847408e+0
1 Alpha  3.355445e+0
6
-----

```

```

./diabetes-5-fold/train/
      Values

```

```

MAE      2.783232e-01
Alpha    1.677722e+0
6
-----

```

```

./mortgage-5-fold/train/
      Values

```

```

MAE      0.005681
Alpha    153.900001
-----

```

```

./plastic-5-fold/train/
      Values

```

```

MAE      2.332295
Alpha    26265.600002
-----

```

```

./stock-5-fold/train/
      Values

```

```

MAE      1.240341
Alpha     5.200002
-----

```

### 0.0.8 Ridge for all using R2 score

```

[33]: def PolynomialRidge(train_file, test_file, header, degree):
      train_df = pd.read_csv(train_file, header=header, delimiter=",")
      test_df = pd.read_csv(test_file, header=header, delimiter=",")

```



```

X_train = train_df.iloc[:, :-1].values
y_train = train_df.iloc[:, -1].values

X_test = test_df.iloc[:, :-1].values
y_test = test_df.iloc[:, -1].values

from sklearn.linear_model import Ridge
from sklearn.preprocessing import PolynomialFeatures

poly_reg=PolynomialFeatures(degree=degree)
X_poly=poly_reg.fit_transform(X_train)

alphas=np.array([2**i for i in range(-18, 30)])

best_alpha, best_mse=None, -float('inf')
for alpha in alphas:
    regressor = Ridge(alpha=alpha)
    regressor.fit(X_poly, y_train)
    y_pred = regressor.predict(poly_reg.transform(X_test))

    from sklearn.metrics import mean_squared_error, mean_absolute_error, _
r2_score
    mse = r2_score(y_test, y_pred)
    if(mse>best_mse):
        best_mse, best_alpha=mse, alpha

[34]: for train_dir, test_dir, header in zip(train_dirs, test_dirs, headers):
    train_files = os.listdir(train_dir)
    test_files = os.listdir(test_dir)

    val = np.zeros(2)
    for train, test in zip(train_files, test_files):
        val += PolynomialRidge(train_dir + train, test_dir + test, header, 2)
    print(train_dir)
    val /= len(train_files)
    val = pd.DataFrame(val, index=["R2", "Alpha"], columns=["Values"])
    print(val)
    print("-----\n")

```

./concrete5-fold/train/

	Values
R2	7.871333e-01
Alpha	3.355445e+0
6	

-----

```
./diabetes-5-fold/train/
      Values
R2      3.330702e-01
Alpha    1.677722e+0
6
-----
```

```
./mortgage-5-fold/train/
      Values
R2      0.999395
Alpha  153.900001
-----
```

```
./plastic-5-fold/train/
      Values
R2      0.799613
Alpha  26265.600002
-----
```

```
./stock-5-fold/train/
      Values
R2      0.970870
Alpha   5.200002
-----
```

```
[35]: for train_dir, test_dir, header in zip(train_dirs, test_dirs, headers):
      train_files = os.listdir(train_dir)
      test_files = os.listdir(test_dir)

      val = np.zeros(2)
      for train, test in zip(train_files, test_files):
          val += PolynomialRidge(train_dir + train, test_dir + test, header, 3)
      print(train_dir)
      val /= len(train_files)
      val = pd.DataFrame(val, index=["R2 score", "Alpha"], columns=["Values"])
      print(val)
      print("-----\n")
```

```
./concrete-5-fold/train/
      Values
R2 score 8.650418e-
01      Alpha
1.074266e+08
-----
```

```
./diabetes-5-fold/train/
      Values
R2 score 1.826710e-
01      Alpha
1.090521e+08
```

```

-----
./mortgage-5-fold/train/
      Values
R2 score 9.995106e-
01      Alpha
1.250165e+08
-----

```

```

-----
./plastic-5-fold/train/
      Values
R2 score
0.81367
1
Alpha 25.625002
-----

```

```

-----
./stock-5-fold/train/
      Values
R2 score 0.983591
Alpha 16896.000001
-----

```